



# Practical Record And Replay Debugging With rr

Robert O'Callahan


# Debugging nondeterminism

Linux opt	B Cpp Jit1 Jit2 Mn Mn-e10s Wr X M( 1 2 3 4 5 JP bc1 bc2 bc3 dt gl oth p) M-e10s( 1 2 3 4 5 bc1 bc2 bc3 dt) R( C J R1 R2 Ru) R-e10s( C R-e10s) T( c d g1 g2 o s tp) W( 1 2 3 4)
Linux pgo	B Cpp Jit1 Jit2 Mn Mn-e10s Wr X M( 1 2 3 4 5 JP bc1 bc2 bc3 dt gl oth p) M-e10s( 1 2 3 4 5 bc1 bc2 bc3 dt) R( C J R1 R2 Ru) R-e10s( C R-e10s) T( c d g1 g2 o s tp) W( 1 2 3 4)
Linux debug	B Cpp Jit1 Jit2 Mn X M( 1 2 3 4 5 JP bc1 bc2 bc3 dt1* dt2 dt3 dt4 gl oth* p) M-e10s( 1 2 3 4 5 bc1* bc1 bc2* bc3) R( C J R1 R2) R-e10s( R-e10s1 R-e10s2)
Linux x64 opt	B Cpp H Jit1 Jit2 Ld Mn V Wr X M( 1 2 3 4 5 JP bc1 bc2 bc3 dt gl oth p) M-e10s( 1 2 3 4 5 bc1 bc2 bc3 dt) R( C J R) R-e10s( C R-e10s) T( c d g1 g2 o s tp) W( 1 2 3 4)
Linux x64 pgo	B Cpp Jit1 Jit2 Ld Mn Wr X M( 1 2 3 4 5 JP bc1 bc2 bc3 dt gl oth p) M-e10s( 1 2 3 4 5 bc1 bc2 bc3 dt) R( C J R) R-e10s( C R-e10s) T( c d g1 g2 o s tp) W( 1 2 3 4)
Linux x64 asan	Bd Bo Cpp Jit1 Jit2 M( 1 2 3 4 5 JP bc1* bc2 bc3 dt* gl oth p) M-e10s( 1 2* 2 3 4 5 bc1* bc2 bc3) R( C J R* )
Linux x64 debug	B Cpp Jit1 Jit2 Mn S X M( 1 2 3 4 5 JP bc1 bc2 bc3 dt1 dt2 dt3 dt4 gl oth p) M-e10s( 1 2 3 4 5 bc1 bc2 bc3) R( C J R1 R2) R-e10s( R-e10s1 R-e10s2)


# Deterministic hardware




# Sources of nondeterminism



# Record inputs



# Replay execution



# “Old idea”

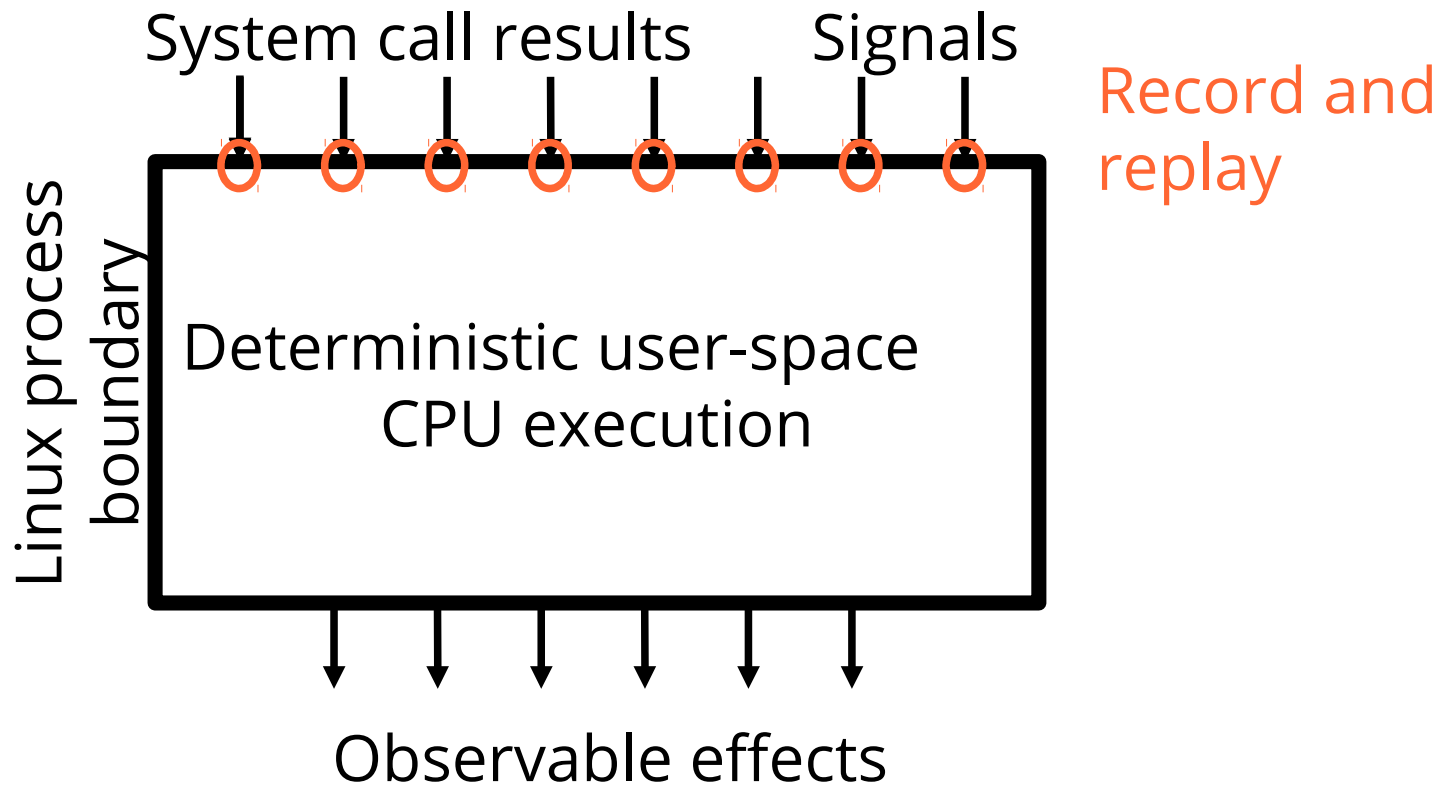
	PinPlay	ReVirt
Nirvana		
	Jockey	ReSpec
Chronomancer		ODR
	PANDA	
	Scribe	Echo
CLAP	FlashBack	QuickRec
ReTrace		

# rr goals

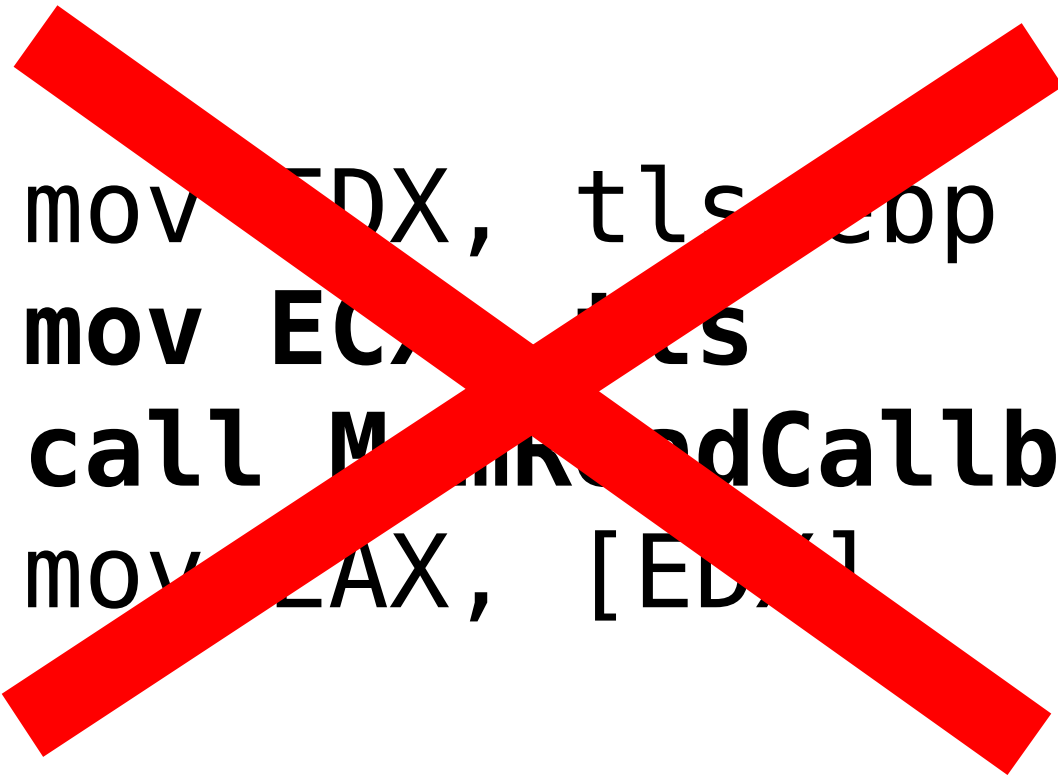
- ▮ Easy to deploy: stock hardware, OS
- ▮ Low overhead
- ▮ Works on Firefox
- ▮ Small investment



# rr design



# No code instrumentation



```
mov EDX, tls_ebp  
mov ECX, tls  
call MarkedCallback  
mov EAX, [EDX]
```

# Use modern HW/OS features

System call results

ptrace

Signals

ptrace

Shared memory data races

Limit to single core

Asynchronous event timing

HW performance counters

Trap on a subset of system calls

seccomp-bpf

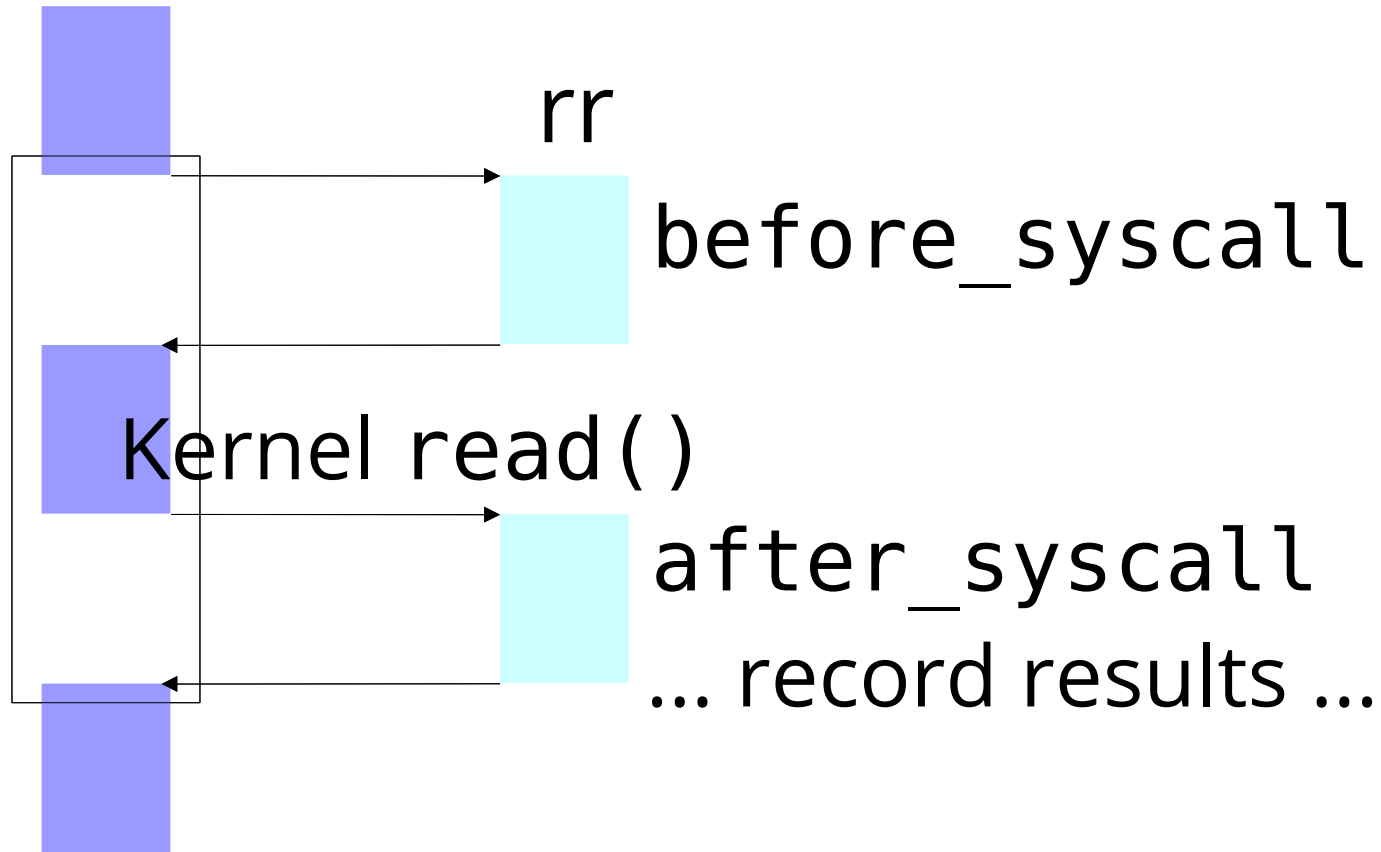
Notification when system call  
blocks in the kernel

DESCHED perf events

Cheap block copies

FIOCLONERANGE

# ptrace



# Use modern HW/OS features

System call results

ptrace

Signals

ptrace

Shared memory data races

Limit to single core

Asynchronous event timing

HW performance counters

Trap on a subset of system calls

seccomp-bpf


Notification when system call  
blocks in the kernel

DESKED perf events

Cheap block copies

FIOCLONERANGE

# Data races



# Data races



# Use modern HW/OS features

System call results

ptrace

Signals

ptrace

Shared memory data races

Limit to single core

Asynchronous event timing

HW performance counters

Trap on a subset of system calls

seccomp-bpf

Notification when system call  
blocks in the kernel


DESKED perf events

Cheap block copies

FIOCLONERANGE



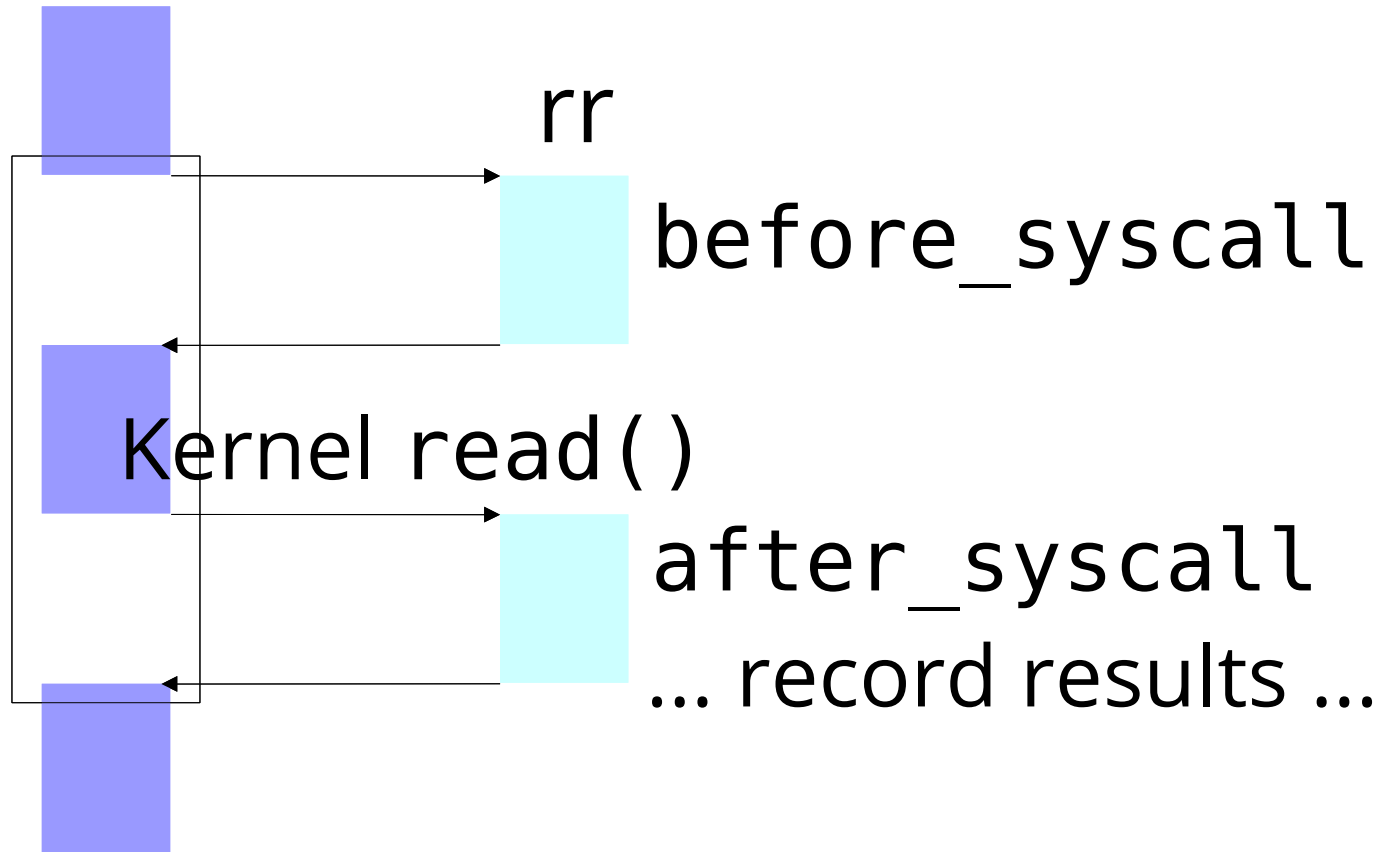
# Event timing: HW perf counters




# Use modern HW/OS features

System call results	ptrace
Signals	ptrace
Shared memory data races	Limit to single core
Asynchronous event timing	HW performance counters
Trap on a subset of system calls	seccomp-bpf
Notification when system call blocks in the kernel	DETSCHED perf events
Cheap block copies	FIOCLONERANGE

# Accelerating system calls

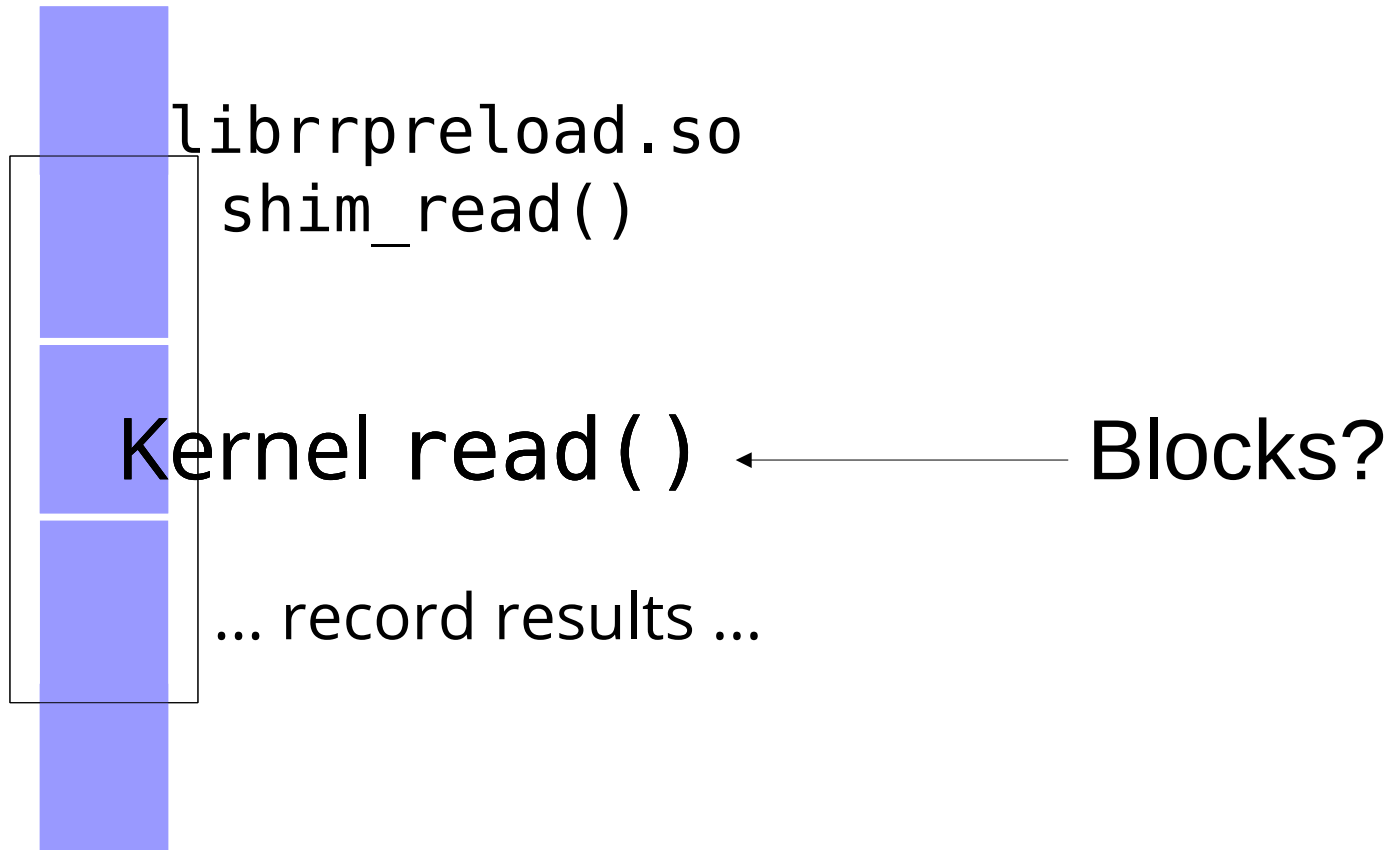


# Avoid context switches



Suppress ptrace trap  
Use seccomp-bpf  
predicates

# Blocking system calls



# Blocking system calls



# Other issues

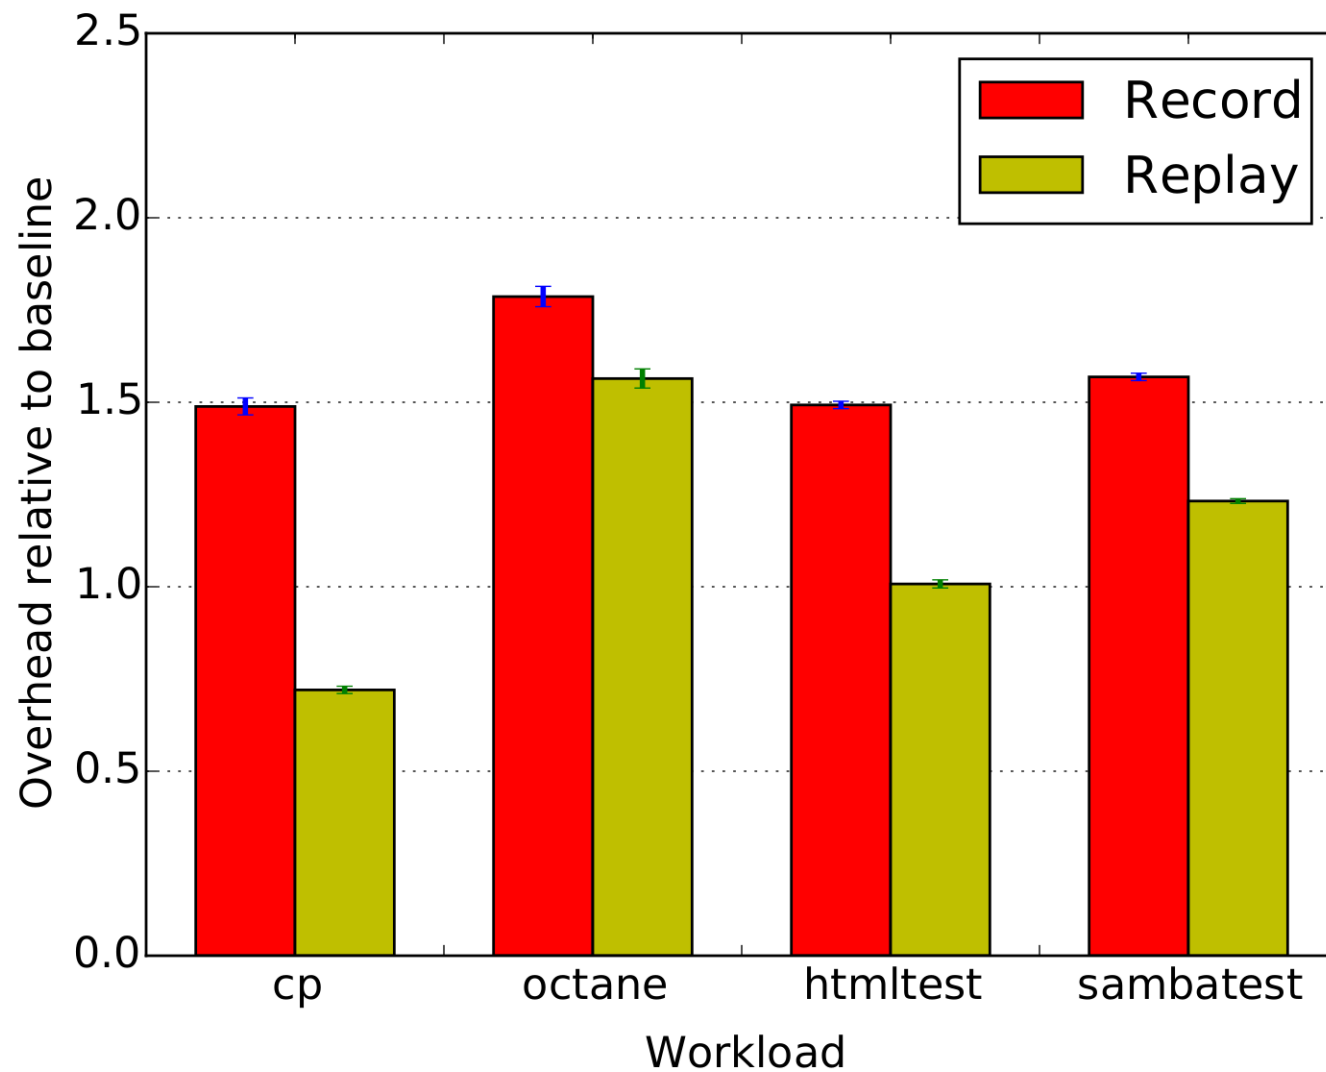
**RDTSC**

**CPUID**

**RDRAND**

**XBEGIN/XEND**

rr Overhead





Debug with gdb

# Running debuggee code

```
(gdb) call ::DumpJSStack()
```

```
0 _setMaxHeight()
```

```
    ["panelUI.xml":331]
```

```
    this = [object XULElement]
```

```
1 handleEvent(aEvent = [object  
MouseEvent])
```

```
    ["panelUI.xml":304]
```

```
    this = [object XULElement]
```

# Running debuggee code

replay



diversion clone



breakpoint

resume



::DumpJSStack()



# Reverse execution

```
(gdb) watch -1 mRect.width
```

```
(gdb) reverse-continue
```

```
nsIFrame::SetRect
```


```
(this=0x2aaadd7db0,
aRect=...)
```

```
718          mRect = aRect;
```

```
(gdb) reverse-next
```

# Reverse execution

replay



# Results

21:38 **mstange** roc: there's somewhat of a competition going on here at the office about who can use rr the most

21:38 **mstange** roc: it's so good

21:39 **roc** :-)

21:39 **roc** who's using it?

21:39 **mstange** roc: jeff, myself, jeff's interns

21:40 **mstange** roc: and we're telling everybody else to use it whenever we get the chance

# Limitations

# Single-core

Recording/replaying inter-core data races

→ need HW support :-(

→ need users, to make economic argument

Find bugs in parallel programs

→ evil scheduler (*chaos mode*)



# ARM

*retry:*

LDREX r0, [addr]

ADD r0, 1

*hardware interrupt???*

STREX r1, r0, [addr]

CMP r1, 0

BNE *retry*

→ Need hardware support to detect/compensate

→ Or binary rewriting

# Kernel semantics

ioctl



Edge cases in system calls


Overhead of switching to supervisor process  
between each tracee context switch


→ Build into OS/hypervisor???

# `gdb`

Not the ultimate debugger interface







```
static uint8_t dr_reg_to_scratch_mask(reg_id_t reg) {  
    if (reg >= DR_REG_R8 && reg <= DR_REG_R11) {  
        return 1 << (reg - DR_REG_R8);  
    }  
    if (reg >= DR_REG_R8D && reg <= DR_REG_R11D) {  
        return 1 << (reg - DR_REG_R8D);  
    }  
    if (reg >= DR_REG_R8W && reg <= DR_REG_R11W) {  
        return 1 << (reg - DR_REG_R8W);  
    }  
    if (reg >= DR_REG_R8L && reg <= DR_REG_R11L) {  
        return 1 << (reg - DR_REG_R8L);  
    }  
    return 0;  
}
```

# Debugging



data analysis and  
visualization!





<http://rr-project.org>  
[mozilla/rr](http://rr-project.org/mozilla/rr)

<https://github.com/>